

# Exploring Regular Expression Usage and Context in Python

Carl Chapman, Kathryn T. Stolee\*

Iowa State University, North Carolina State University

*carlallenchapman@gmail.com, ktstolee@ncsu.edu*

19 July, 2016

# Why regular expressions?

# Why regular expressions?

- Regexes are everywhere! (we think...)

# Why regular expressions?

- Regexes are everywhere! (we think...)
- Everyone writes regexes! (we think...)

# Why regular expressions?

- Regexes are everywhere! (we think...)
- Everyone writes regexes! (we think...)
- Regexes are hard to read/write! (again, we think...)

# Why regular expressions?

- So, we wanted to write a tool to support regex creation.

# Why regular expressions?

- So, we wanted to write a tool to support regex creation.
- But...

# Why regular expressions?

- So, we wanted to write a tool to support regex creation.
- But...

Regex feature usage references are missing!



# Why regular expressions?

- So, we wanted to write a tool to support regex creation.
- But...

Regex feature usage references are missing!

- and...

# Why regular expressions?

- So, we wanted to write a tool to support regex creation.
- But...

Regex feature usage references are missing!

- and...

We don't know how/when/why developers use regexes!

# Research goals

## Explore regex

- 1 Context (developer survey)
- 2 Features (repository analysis)
- 3 Use cases (similarity analysis)

# Regular expressions: The basics

- `(ab*c|yz*)$`

✓ abbbbbbbbc

✓ y

✓ abcy

# Regular expressions: The basics

- `(ab*c|yz*)$`

✓ abbbbbbbbc

✓ y

✓ abcy

✗ abcccc

✗ yxw

# Regular expressions: The basics

- `(ab*c|yz*)$`

✓ abbbbbbbbc

✓ y

✓ abcy

✗ abcccc

✗ yxw

- `(ab*c|yz*)`

# Regular expressions: The basics

- `(ab*c|yz*)$`

✓ abbbbbbbbc

✓ y

✓ abcy

✗ abcccc

✗ yxw

- `(ab*c|yz*)`

✓ abbbbbbbbc

✓ y

✓ abcy

✓ abcccc

✓ yxw

# Part 1

## RQ1

In what contexts do professional developers use regular expressions?



# Survey context

- 18 professional developers
- 9 years average development experience
- Small mobile payment management company
- 30 questions in a Google form

# How frequently do developers use regexes?

- 50% – at least once per week
- Regexes are most frequently composed within command line and text editor tools
- 2 developers write more than 50 regexes in general programming languages (e.g., Java) annually
- Database queries using regexes were rare

# Common regex activities

How often do you use regexes for...

Activity	Frequency
Locating content within a file or files	4.4
Capturing parts of strings	4.3
Parsing user input	4.0

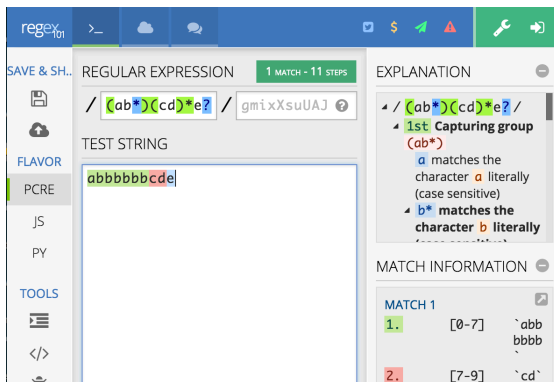
Key: 6 = very frequently, 5 = frequently, 4 = occasionally,  
3 = rarely, 2 = very rarely, 1 = never

# Testing regular expressions

Developers test regular expressions less often than other code.

# Testing regular expressions

Developers test regular expressions less often than other code.



50% say they use testing tools like [www.regex101.com](http://www.regex101.com)

# Pain points

hard to compose (11 = 61%)

...very difficult to write them since I've never read up on them.

...trickiness to getting the expression right

# Pain points

hard to compose (11 = 61%)

...very difficult to write them since I've never read up on them.

...trickiness to getting the expression right

hard to read (7 = 39%)

long ones can be hard to read

Readability. Edge cases.

It is terrible to read (especially later after initial development)

# Pain points

hard to compose (11 = 61%)

...very difficult to write them since I've never read up on them.

...trickiness to getting the expression right

hard to read (7 = 39%)

long ones can be hard to read

Readability. Edge cases.

It is terrible to read (especially later after initial development)

inconsistency across implementations (3 = 17%)

Differences in implementation across languages

Some regexes work differently (or don't work) in some languages.



# Notable Observations

- Regexes are composed fairly frequently by developers
- Testing regexes is less common than testing other code
- Developers find regexes hard to read and write

# Notable Observations

- Regexes are composed fairly frequently by developers
- Testing regexes is less common than testing other code
- Developers find regexes hard to read and write

but....

# Notable Observations

- Regexes are composed fairly frequently by developers
- Testing regexes is less common than testing other code
- Developers find regexes hard to read and write

but....



How do developers really use regexes?

## Part 2

### RQ3

Which regular expression language features are most commonly used in Python?

# Project selection with the GitHub API



Out of 3,898 pseudo-randomly selected Python projects, 1,645 (42%) contained one or more regex utilization.

# In Python: Utilizations of the re module

```
function      pattern      flags  
r1 = re.compile("(0|-?[1-9][0-9]*)$", re.MULTILINE)
```

**function** which function of the re module is called?

**pattern** string used to specify regex behavior

**flags** modifies the regex engine

# Filtering utilizations and patterns

**53,894** unique utilizations observed.

12.7% use behavioral flags

6.5% were non-static patterns

**43,525** utilizations remain

**13,711** distinct normalized patterns

114 had various errors

**13,597** usable patterns remain for analysis

# Filtering utilizations and patterns

**53,894** unique utilizations observed.

12.7% use behavioral flags

6.5% were non-static patterns

**43,525** utilizations remain

**13,711** distinct normalized patterns

114 had various errors

**13,597** usable patterns remain for analysis

Average utilizations per project: 32



# PCRE Parsing Patterns



# PCRE Parsing Patterns



All Python features are recognizable by PCRE

# Feature statistics - Top 8

Rank	Code	Example	% Projects	NProjects	NFiles	NPatterns	MaxTokens
1	ADD	z+	73.2	1,204	9,165	6,003	30
2	CG	(caught)	72.6	1,194	9,559	7,130	17
3	KLE	.*	66.8	1,099	8,163	6,017	50
4	CCC	[aeiou]	62.4	1,026	7,648	4,468	42
5	ANY	.	61.1	1,005	6,277	4,657	60
6	RNG	[a-z]	51.6	848	5,092	2,631	50
7	STR	^	51.4	846	5,458	3,563	12
8	END	\$	50.3	827	5,393	3,169	12

# Regex research tools

- Remember that we wanted to write a tool to support regex creation?

# Regex research tools

- Remember that we wanted to write a tool to support regex creation?
- But regex feature usage references were missing.

# Regex research tools

- Remember that we wanted to write a tool to support regex creation?
- But regex feature usage references were missing.
- So,

# Regex research tools

- Remember that we wanted to write a tool to support regex creation?
- But regex feature usage references were missing.
- So,

How well do past and present regex research tools meet the needs of developers? (Hampi, Rex, RE2, brics)

# Which features are supported by analysis tools?

Rank	Code	Example	Brics	Hampi	Rex	RE2
1	ADD	z+	●	●	●	●
2	CG	(caught)	●	●	●	●
3	KLE	.*	●	●	●	●
4	CCC	[aeiou]	●	●	●	●
5	ANY	.	●	●	●	●
6	RNG	[a-z]	●	●	●	●
7	STR	^	●	●	●	●
8	END	\$	●	●	●	●
9	NCCC	[^qwx]	●	●	●	●
10	WSP	\s	●	●	●	●
11	OR	a b	●	●	●	●
12	DEC	\d	●	●	●	●
13	WRD	\w	●	●	●	●
14	QST	z?	●	●	●	●
15	LZY	z+?	●	●	●	●
16	NCG	a(?:b)c	●	●	●	●
17	PNG	(?P<name>x)	●	●	●	●

Rank	Code	Example	Brics	Hampi	Rex	RE2
18	SNG	z{8}	●	●	●	●
19	NWSP	\s	●	●	●	●
20	DBB	z{3,8}	●	●	●	●
21	NLKA	a(?:yz)	●	●	●	●
22	WNW	\b	●	●	●	●
23	NWRD	\w	●	●	●	●
24	LWB	z{15,}	●	●	●	●
25	LKA	a(?:bc)	●	●	●	●
26	OPT	(?i)CasE	●	●	●	●
27	NLKB	(?!x)yz	●	●	●	●
28	LKB	(?<=a)bc	●	●	●	●
29	ENDZ	\Z	●	●	●	●
30	BKR	\1	●	●	●	●
31	NDEC	\D	●	●	●	●
32	BKRN	\g<name>	●	●	●	●
33	VWSP	\v	●	●	●	●
34	NWNW	\B	●	●	●	●



# Notable Observations

- Regexes are common, but not everywhere (42% of projects, 32 utilizations per project)
- Current regex research tools cover the most common features

# Notable Observations

- Regexes are common, but not everywhere (42% of projects, 32 utilizations per project)
- Current regex research tools cover the most common features

but....

# Notable Observations

- Regexes are common, but not everywhere (42% of projects, 32 utilizations per project)
- Current regex research tools cover the most common features

but....

What are the regexes doing?

## Part 3

### RQ4

How behaviorally similar are regexes across projects?

# How to find common behaviors?

- 1 ~~thorough inspection of 53K utilizations~~
- 2 ~~cluster by syntactic similarity like Jaccard or longest substring~~
- 3 ~~formal analytical subsumption, no sufficient tools at the moment~~
- 4 Chosen technique: cluster by behavioral similarity using Rex

# Example

A (ab\*c|yz\*)\$

- abbbbbbbbc
- y
- abcy
- pac
- abcyzzz

B (ab\*c|yz\*)

- y
- abc
- abcy
- abcccc
- yxw

# Example

A (ab\*c|yz\*)\$

- abbbbbbbbc
- y
- abcy
- pac
- abcyzzz

B (ab\*c|yz\*)

- y
- abc
- abcy
- abcccc
- yxw

A matches 3/5 =  
60% of B's strings

# Example

A (ab\*c|yz\*)\$

- abbbbbbbbc
- y
- abcy
- pac
- abcyzzz

A matches 3/5 =  
60% of B's strings

B (ab\*c|yz\*)

- y
- abc
- abcy
- abcccc
- yxw

B matches 5/5 =  
100% of A's strings



# Example

A (ab\*c|yz\*)\$

- abbbbbbbbc
- y
- abcy
- pac
- abcyzzz

A matches 3/5 =  
60% of B's strings

B (ab\*c|yz\*)


- y
- abc
- abcy
- abcccc
- yxw

B matches 5/5 =  
100% of A's strings

A and B are 80% similar

# Similarity Matrix → MCL

	A	B	C	D
A	1.0	0.0	0.9	0.0
B	0.2	1.0	0.8	0.7
C	0.6	0.8	1.0	0.2
D	0.0	0.6	0.1	1.0



	A	B	C	D
A	1.0			
B	0.1	1.0		
C	0.75	0.8	1.0	
D	0.0	0.65	0.15	1.0

Rex generates  
400 strings for each regex.  
Average scores to  
half-matrix for MCL

# Scope

- 3,582 (26%) of patterns appeared in multiple projects
- 711 unsupported by Rex

# Scope

- 3,582 (26%) of patterns appeared in multiple projects
- 711 unsupported by Rex
- 2,871 patterns analyzed from 722 (44%) of the projects

# Clustering Results

From 2,871 distinct regexes:  
→ 186 clusters with size  $\geq 2$   
→ 2,042 unclustered regexes

# Clustering Results

## Example Cluster

Index	Pattern	NProjects	Index	Pattern	NProjects
1	<code>\s*([[: ]*)\s*:(.*)</code>	9	7	<code>[:]</code>	6
2	<code>:+</code>	8	8	<code>([[: ]+):(.*)</code>	6
3	<code>(:)</code>	8	9	<code>\s*:\s*</code>	4
4	<code>(: +)</code>	8	10	<code>\:</code>	2
5	<code>(:)(:*)</code>	8	11	<code>^([[: ]*)[:[: ])*\$</code>	2
6	<code>^([[: ]*)[: *(.*)</code>	8	12	<code>^([[: ]*)[:([[: ]*)\$</code>	2

From 2,871 distinct regexes:  
 → 186 clusters with size  $\geq 2$   
 → 2,042 unclustered regexes

# Six Categories Of Clusters

Category	Clusters	Patterns	Projects	% Projects
Multi Matches	21	237	295	40%
Specific Char	17	103	184	25%
Anchored Patterns	20	85	141	19%
Two or More Chars	16	40	120	16%
Content of Parens	10	46	111	15%
Code Search	15	27	92	13%

Multi Matches `(\s)`, `,`, `|`;

Specific Char `:+`, `}`, `%`

Anchored Patterns

`^[-_A-Za-z0-9]+$`

Two Or More Chars `@[a-z] +`

Content of Parens `<(.)>`,

`<[>]*?>`

Code Search `.*rlen=([0-9]+)`

# Notable Observations

- Finding a specific character is quite common, 25% of projects (in contrast with survey)
- Regexes are often used to parse source code!



# Notable Observations

- Finding a specific character is quite common, 25% of projects (in contrast with survey)
- Regexes are often used to parse source code!

but....

# Notable Observations

- Finding a specific character is quite common, 25% of projects (in contrast with survey)
- Regexes are often used to parse source code!

but....

- Similarity metric is approximate
- Metric is perhaps too sensitive to differences in literals

# Opportunities for future work!

# Opportunities for future work!

## Refactoring Regular Expressions

Regexes are hard to read. Could refactoring help?

# Opportunities for future work!

## Refactoring Regular Expressions

Regexes are hard to read. Could refactoring help?

## Migration Support for Developers

Supported regex features are different among languages.

# Opportunities for future work!

## Refactoring Regular Expressions

Regexes are hard to read. Could refactoring help?

## Migration Support for Developers

Supported regex features are different among languages.

## Better Similarity Metrics

Our similarity metrics are empirical, can we do it analytically?

# Opportunities for future work!

## Refactoring Regular Expressions

Regexes are hard to read. Could refactoring help?

## Migration Support for Developers

Supported regex features are different among languages.

## Better Similarity Metrics

Our similarity metrics are empirical, can we do it analytically?

## Identifying Best Practices?

Could impact regex education and improve comprehension.

# Recap

- Regexes are (sort of) everywhere! (42% of Python projects studied, but more frequently in text editors and IDEs)



# Recap

- Regexes are (sort of) everywhere! (42% of Python projects studied, but more frequently in text editors and IDEs)
- Everyone writes regexes! (50% of devs write them weekly)

# Recap

- Regexes are (sort of) everywhere! (42% of Python projects studied, but more frequently in text editors and IDEs)
- Everyone writes regexes! (50% of devs write them weekly)
- Regexes are hard to read/write! (this is a pain point)

# Recap

- Regexes are (sort of) everywhere! (42% of Python projects studied, but more frequently in text editors and IDEs)
- Everyone writes regexes! (50% of devs write them weekly)
- Regexes are hard to read/write! (this is a pain point)

also...

- Current tools support most of the most common features
- Regexes are often used for code search
- Many opportunities for future work!

# Questions?

Katie Stolee – ktstolee@ncsu.edu

# Questions?

Katie Stolee – ktstolee@ncsu.edu

(psst! Graduate students! I'm hiring!)